

Cassandra Cluster Analysis.

Dongyao Wang (dowa4431@colorado.edu), Sandesh Dhawaskar Sathyanarayana (sadh0344@colorado.edu), Sananda Banerjee(Sananda.Banerjee@colorado.edu)

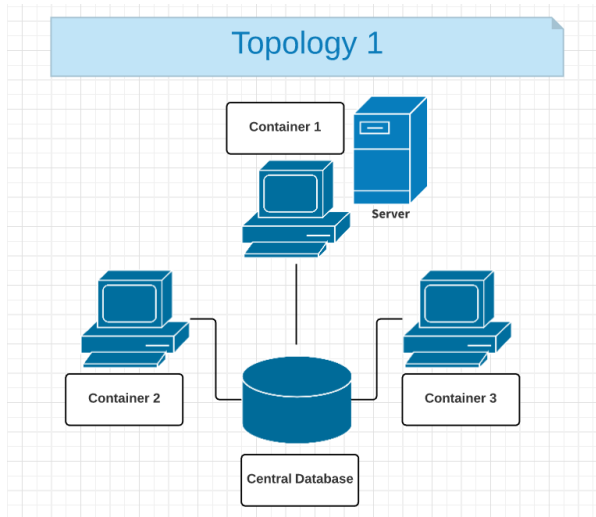
1. Introduction

Clock Synchronization is one of the major challenge and issue in the Distributed Data Bases like Cassandra. NTP is one of the main tool used for the clock synchronization. NTP measure the clock difference and use the same to either decrease or increase the clock speed in the individual node/machine. This project has two major phases with two major analysis. In the first phase we look at the analysis of the delay, offset of the clock with NTP servers and in the second phase we measure the read/write latencies of the distributed data base. Cassandra is a key-value store distributed database, essentially a giant hash table. Hence, we choose the Cassandra for our analysis. We also aim at identifying and developing tools and methodologies to facilitate running and testing of Cassandra Cluster, across various topologies and network conditions. We gained experience with the Apache Cassandra database. There are three main topologies that have been implemented and we will look at each one of them in detail in the following sections.

2. Topology 1

2.1 Description – setup, configuration

For the first topology we had to setup three



Fig(1)

machines on the same device. So, we created three containers called Cassandra nodes and named them cn1, cn2 and cn3 as

can be seen in the figure Fig (1). We established this arrangement using docker on the virtual machine of the same computer.

To test if the three instances are up and running, we used the command `node tool status` that shows the status of the machines (containers). Moreover, we opened three separate tabs for the three containers and tried pinging between them. We used the command `ifconfig` to determine the IP address of the current machine and accordingly pinged to the other two nodes. The nodes could successfully ping each other, sending and receiving packets. This confirms that Cassandra cluster has formed.

The below snapshot shows that three containers are up and running.

```
File Edit View Search Terminal Help
root@sandesh-VirtualBox:~# docker ps -a
CONTAINER ID        IMAGE               COMMAND             CREATED             STATUS              PORTS
86f1fdaa2436        cassandra:latest   "docker-entrypoint.s..." 3 weeks ago         Up About a minute   7000-7001/tcp, 7199/tcp, 9042/tcp, 91
60/tcp
68ec36f375a        cassandra:latest   "docker-entrypoint.s..." 3 weeks ago         Up About a minute   7000-7001/tcp, 7199/tcp, 9042/tcp, 91
60/tcp
dba8229523c        cassandra:latest   "docker-entrypoint.s..." 3 weeks ago         Up About a minute   7000-7001/tcp, 7199/tcp, 9160/tcp, 0.
0.0.0:9042->9042/tcp
root@sandesh-VirtualBox:~#
```

After this you should login to each container and see if each of them has got different IP address. Below snap shots give that information.

```
root@dba8229523c:~# ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:ac:12:00:02
          inet addr:172.18.0.2  Bcast:172.18.255.255  Mask:255.255.0.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:6096 errors:0 dropped:0 overruns:0 frame:0
          TX packets:5979 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:672055 (656.3 KiB)  TX bytes:658714 (643.2 KiB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:24 errors:0 dropped:0 overruns:0 frame:0
          TX packets:24 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:1370 (1.3 KiB)  TX bytes:1370 (1.3 KiB)

root@dba8229523c:~#
```

```
root@68ec36f375a:~# ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:ac:12:00:03
          inet addr:172.18.0.3  Bcast:172.18.255.255  Mask:255.255.0.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:6413 errors:0 dropped:0 overruns:0 frame:0
          TX packets:6313 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:706057 (689.5 KiB)  TX bytes:690267 (674.0 KiB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:24 errors:0 dropped:0 overruns:0 frame:0
          TX packets:24 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:1370 (1.3 KiB)  TX bytes:1370 (1.3 KiB)

root@68ec36f375a:~#
```

```
root@68ec36f375a:~# ifconfig
eth0      Link encap:Ethernet  HWaddr 02:42:ac:12:00:03
          inet addr:172.18.0.3  Bcast:172.18.255.255  Mask:255.255.0.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:6413 errors:0 dropped:0 overruns:0 frame:0
          TX packets:6313 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:706057 (689.5 KiB)  TX bytes:690267 (674.0 KiB)

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  Metric:1
          RX packets:24 errors:0 dropped:0 overruns:0 frame:0
          TX packets:24 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:1370 (1.3 KiB)  TX bytes:1370 (1.3 KiB)

root@68ec36f375a:~#
```

After this you should login to machine one of them and check if cluster has formed. Also, they will show the different IP's we

showed above.

```
File Edit View Search Terminal Help
root@86f1fdaa2436:~# nodetool status
datacenter1 datacenter1
=====
Status=Up/Down
/- State=Normal/Leaving/Joining/Moving
-- Address          Load          Tokens         Owns (effective)  Host ID                               Rack
jn 172.18.0.2        3.63 MiB      250             70.7%              c7f0ef60-f368-4418-aae2-b2201a676cfe  rack1
jn 172.18.0.3        3.4 MiB       250             66.4%              7a5df6c0-78af-4965-bf8e-4e47c213a738  rack1
jn 172.18.0.4        3.25 MiB     250             62.5%              7d681e50-2990-4d3c-ac53-8395677006fa  rack1

root@86f1fdaa2436:~#
```

Next, we took the testing it a step forward, by exploring the Cassandra Query language, also known as the CQL. Inside cqlsh of one node (say cn1), we created a key space and inside the key space we created tables where we inserted and viewed values. Then, we looked up key spaces from a second node (say cn2) and we could view the same table from cn2 as well. We could insert and update values on the table and the changes would be reflected in the table and could be verified from the other nodes. Similarly, we tested with cn3 and obtained similar results. Following snapshots will describe the example explained above.

```
root@sandesh-VirtualBox:~# cqlsh> desc KEYSPACES
system_schema system system_distributed sandesh
system_auth "TestKEYSPACE1" system_traces cn1_cn2_cn3
cqlsh> desc KEYSPACES ;
cqlsh> use cn1_cn2_cn3 ;
cqlsh> desc TABLES ;
empty>
cqlsh:cn1_cn2_cn3> desc tables;
table1
cqlsh:cn1_cn2_cn3> insert into table1(id)values(20);
cqlsh:cn1_cn2_cn3>
cqlsh:cn1_cn2_cn3> insert into table1(id)values(30);
cqlsh:cn1_cn2_cn3>
```


```
root@sandesh-VirtualBox:~# cqlsh:cn1_cn2_cn3> desc tables;
empty>
cqlsh:cn1_cn2_cn3> desc tables;
table1
cqlsh:cn1_cn2_cn3> select * from table1;
....
10
(1 row)
cqlsh:cn1_cn2_cn3> select * from table1;
10
....
20
(1 row)
cqlsh:cn1_cn2_cn3>
```

After this, we used the Cassandra stress tool to generate statistics to measure the latencies for read and write operations as can be viewed in the table below. We obtained interesting results, for various iterations of n. As we incremented the number of iterations, the latency would increase both for read and write as can be seen from the grey dots in the graph below.

2.2. Table:

| Iterations | Read Latency | Write Latency | Write op-Rate(ms) |
|------------|--------------|---------------|-------------------|
| 10 | 28.54 | 35.9 | 141 |
| 50 | 26.28 | 109.3 | 293 |
| 100 | 46.73 | 194 | 429 |
| 150 | 62.8 | 218.3 | 444 |
| 200 | 68.68 | 342.7 | 317 |
| 250 | 42.75 | 281.1 | 501 |
| 300 | 81.13 | 264.3 | 597 |
| 350 | 90 | 270.9 | 565 |
| 400 | 75 | 320.2 | 482 |
| 450 | 97 | 295.1 | 561 |
| 500 | 83.17 | 344.3 | 631 |
| 550 | 78.04 | 233.1 | 565 |
| 600 | 70 | 382.5 | 420 |
| 650 | 78 | 261 | 653 |
| 700 | 68.76 | 255.5 | 452 |
| 750 | 88.13 | 325 | 503 |
| 800 | 92 | 190 | 477 |
| 850 | 44 | 267 | 182 |
| 900 | 64 | 220.6 | 679 |
| 950 | 57 | 206.9 | 624 |
| 1000 | 67 | 251.3 | 656 |
| 1100 | 58 | 286.3 | 618 |
| 1200 | 26 | 215.6 | 811 |
| 1300 | 49 | 229.6 | 797 |
| 1400 | 32 | 224.4 | 750 |
| 1500 | 40 | 230.7 | 727 |
| 1600 | 44 | 211.7 | 893 |
| 1700 | 37.42 | 201.7 | 887 |
| 1800 | 36 | 173.2 | 1074 |
| 1900 | 54 | 166.5 | 1009 |
| 2000 | 43 | 159.2 | 1054 |
| 2500 | 34.15 | 162.8 | 1179 |
| 3000 | 28.46 | 168.2 | 1069 |
| 3500 | 39 | 213.8 | 907 |
| 4000 | 33 | 163.8 | 1125 |
| 4500 | 36 | 159.7 | 1229 |
| 5000 | 34 | 128.4 | 1510 |
| 6000 | 26.25 | 132.9 | 1444 |
| 7000 | 31 | 127.6 | 1470 |
| 8000 | 36 | 120.3 | 1561 |
| 9000 | 33 | 122.4 | 1556 |
| 10000 | 32 | 127.8 | 1210 |
| 20000 | 38 | 106.3 | 1866 |
| 30000 | 37 | 96.4 | 2059 |
| 40000 | 31 | 88.9 | 2236 |
| 50000 | 31.23 | 87.8 | 2267 |
| 60000 | 33.2 | 87.9 | 2269 |
| 70000 | 34.5 | 82.9 | 2404 |
| 80000 | 31.33 | 88.8 | 2245 |
| 90000 | 33.7 | 63.6 | 3126 |
| 100000 | 31.62 | 94.7 | 2107 |

2.3. Graph:



2.4 Observations:

As can be observed from the graph plotting and the values of read and write mean latencies in the table, the op rate keeps increasing linearly with increase in the number of iterations, as can be seen in the grey dots on the graph above. Another major observation being done here is, since all the three containers sit on one Linux machine, no matter how you change the date and time of each node they will always get set to underlying Linux machine date and time.

2.5 Technical Challenges:

- Formation of the cluster with different IP's for each container
- NTP configuration of each node to make cn1 as master node and other two as slaves
- Figuring out how to login to cqlsh with the given IP.

Following snapshots shows the NTP configurations that was used in this

topology. CN1 is the server and cn2, cn3 are the clients.

```

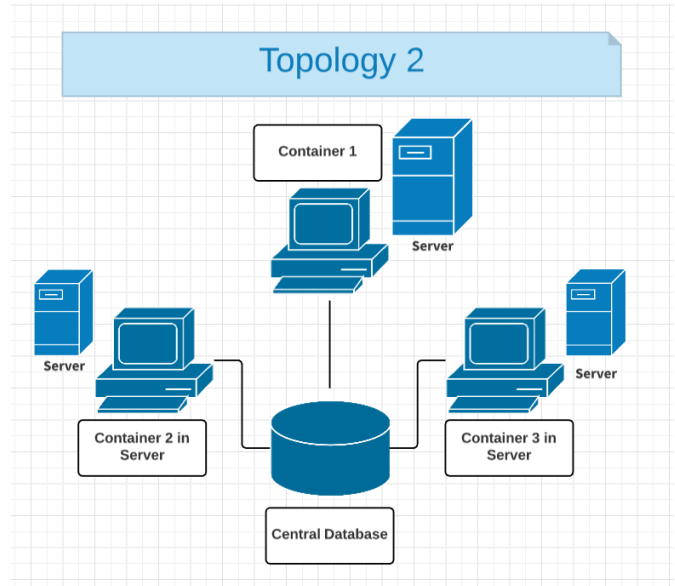
root@dbae8229523c:/# ntpq -p
remote      refid      st t when poll reach  delay  offset jitter
=====
*LOCAL(0)    .LOCL      5 l  63  64  377  0.000  0.000  0.000
root@dbae8229523c:/#

root@86f1fdaa2436:/# ntpq -p
remote      refid      st t when poll reach  delay  offset jitter
=====
*cn1.cluster_net LOCAL(0)  6 u  66  64  377  0.073  0.021  0.016
root@86f1fdaa2436:/#

root@86ec036f375a:/# ntpq -p
remote      refid      st t when poll reach  delay  offset jitter
=====
*cn1.cluster_net LOCAL(0)  6 u  56  64  377  0.068  0.020  0.021
root@86ec036f375a:/#

```

and the other two nodes in the other two VM instances as NTP clients or slave nodes.



3. Topology 2

3.1. Description – setup, configuration

For the second topology, all the Cassandra nodes are locally on the same host, but in separate virtual machines and a single NTP server. In this case each VM runs a single instance of a Cassandra node, using docker containers. So, now we setup two individual containers using docker on the different instances of VM and formed a cluster with a single NTP server.

Then, using commands like nodetool status we tested the current state of all the nodes across the two instances of VM. Then tried pinging between them, sending and receiving packets. Once ping was successful, we tested the status using ntpq, by setting on node as the NTP server or master node,

Following this, we carried out the Cassandra stress test, measuring read write latency values for incrementing number of iterations. We started testing for $n = 10$, where n is the number of iterations and went on until $n = 100000$. The result of the observations of the measurements is discussed in the subsequent sections.

The following snapshot gives the IP's of each host and show that containers are running.

```

root@dbae8229523c:/# ip netns exec cn1 namespace
192.168.1.100
root@dbae8229523c:/# ip netns exec cn2 namespace
192.168.1.101
root@dbae8229523c:/# ip netns exec cn3 namespace
192.168.1.102
root@dbae8229523c:/#

```


3.2. Tables:

- Write


| Iteration | Read op rate(op/s) | Read mean latency(ms) |
|-----------|--------------------|-----------------------|
| 10 | 12 | 122.4 |
| 50 | 46 | 109.4 |
| 100 | 119 | 235.2 |
| 150 | 201 | 133.5 |
| 200 | 170 | 159 |
| 250 | 180 | 418.6 |
| 300 | 472 | 606.4 |
| 350 | 247 | 243 |
| 400 | 293 | 93 |
| 450 | 335 | 371.8 |
| 500 | 381 | 135.8 |
| 550 | 441 | 436.3 |
| 600 | 385 | 82.8 |
| 650 | 523 | 157.3 |
| 700 | 627 | 177.8 |
| 750 | 794 | 134 |
| 800 | 688 | 193.2 |
| 850 | 619 | 188.2 |
| 900 | 699 | 680.4 |
| 950 | 498 | 103.7 |
| 1000 | 419 | 264.8 |
| 1500 | 656 | 133.5 |
| 2000 | 1284 | 909.4 |
| 2500 | 715 | 466.1 |
| 3000 | 1548 | 531.8 |
| 3500 | 1350 | 492.9 |
| 4000 | 1255 | 703.8 |
| 4500 | 1091 | 678.6 |
| 5000 | 825 | 909.1 |
| 5500 | 1171 | 735 |
| 6000 | 1427 | 500.7 |
| 6500 | 1271 | 550 |
| 7000 | 1470 | 625 |
| 7500 | 1090 | 505.4 |
| 100000 | 1681 | 532.6 |

-Read


| Iterations | Write op rate(op/s) | Write mean latency(ms) |
|------------|---------------------|------------------------|
| 10 | 13 | 93.5 |
| 50 | 76 | 215.6 |
| 100 | 94 | 363.1 |
| 150 | 100 | 503.7 |
| 200 | 146 | 551.6 |
| 250 | 140 | 770.2 |
| 300 | 278 | 492.6 |
| 350 | 196 | 433.9 |
| 400 | 216 | 454.2 |
| 450 | 210 | 437.1 |
| 500 | 211 | 482.2 |
| 550 | 262 | 404.4 |
| 600 | 255 | 448.3 |
| 650 | 387 | 416.8 |
| 700 | 421 | 342.5 |
| 750 | 357 | 433.1 |
| 800 | 294 | 399.1 |
| 850 | 309 | 381.1 |
| 900 | 396 | 354.2 |
| 950 | 264 | 447.9 |
| 1000 | 424 | 344.1 |
| 1500 | 368 | 426.3 |
| 2000 | 467 | 326.2 |
| 2500 | 517 | 368.8 |
| 3000 | 504 | 334.5 |
| 3500 | 549 | 320.8 |
| 4000 | 549 | 314.5 |
| 4500 | 762 | 227.7 |
| 5000 | 749 | 257.9 |
| 5500 | 731 | 247 |
| 6000 | 723 | 255.6 |
| 6500 | 706 | 255.4 |
| 7000 | 841 | 206.4 |
| 7500 | 804 | 237.3 |
| 8000 | 662 | 280.8 |
| 8500 | 919 | 201.1 |
| 9000 | 946 | 210.2 |
| 9500 | 907 | 211.8 |
| 10000 | 976 | 198.4 |
| 50000 | 1077 | 183 |
| 100000 | 1277 | 154.9 |

3.3. Graphs:

-Write



-Read



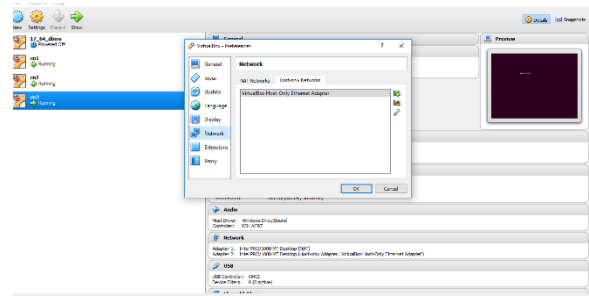
3.4 Observations:


As can be observed from the tables and graphs in the previous sections, the curve rises exponentially with the increase in the number of iterations, much like the first case (topology 1). However, it is interesting to observe that. Second major observation that could be done is regarding the NTP. Now each docker is one each of the VM and hence their time and date would be set on basis of underlying Linux. Hence, we formed underlying Linux kernel as the NTP servers of first node and other two VM kernel machine as the clients of the master for NTP synchronization.

3.5 Technical challenge:


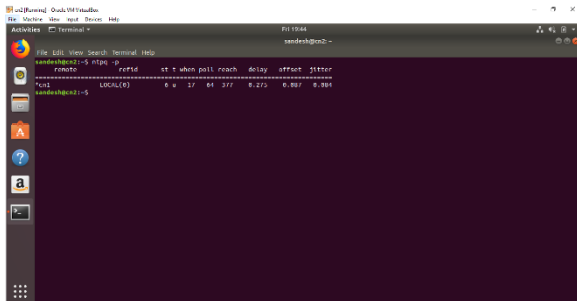

- Formation of the cluster is very difficult task when we have machine one different VM.
- We had to make VM to use host-only machine adaptor and then we had to assign the static IP to each of the machine.
- Once the IP was set we had to form the cluster using the docker.
- Then pinging each machine and forming the cluster.
- Formation of cluster is much harder than the one in topology 1.

Below snapshots give the idea how we achieved the configurations.





Below are some of the NTP configurations snapshots which shows that cn1 was the master and cn2, cn3 were the slaves.

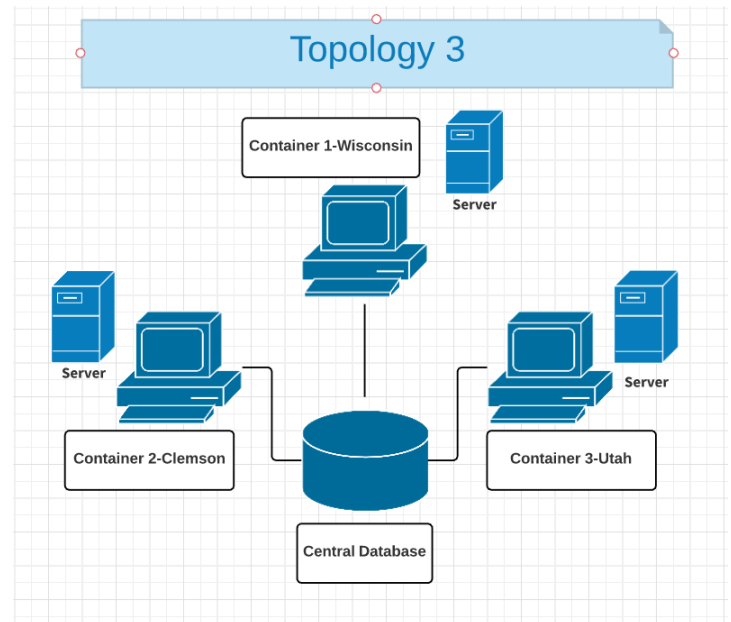


We have not added any database result snapshots as they are exactly as similar to one attached in topology 1.

4. Topology 3

4.1 Description – setup, configuration

For the third topology, we are required to create a cluster with geographically distant Cassandra nodes. To establish this, we used the CloudLab to set up nodes in the cloud and used our device as physical hosts (with nodes on instances of VM) to connect to the nodes in the cloud. As can be seen in the diagram below, all the three nodes are geographically distant from one another, at Wisconsin, Clemson and Utah respectively.



Like the previous two topologies, here also we used node tool status to check the current state of the machine. Then using Network Time Protocol, we set a node as the server and the other two as clients, using the Python codes for server and client, adjusting the IPs to establish the network.

After this, we used Cassandra stress test to overload the network with read and write operations and took measurements for various iterations of read and write operations.

4.2. Tables:

- Write

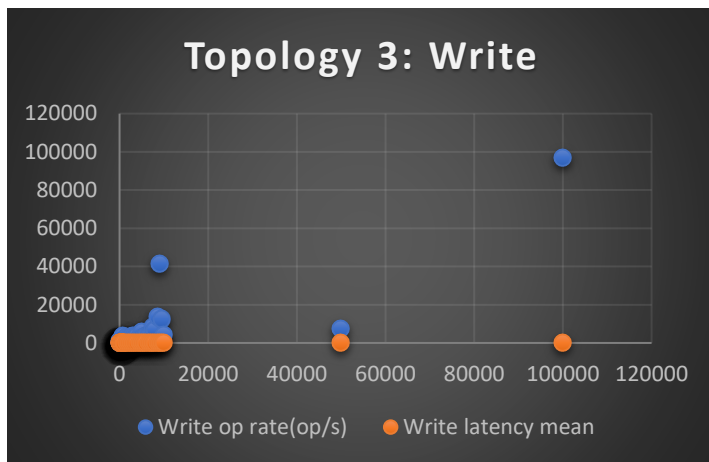
| Iterations | Write op rate(op/s) | Write latency mean |
|------------|---------------------|--------------------|
| 10 | 78 | 12.1 |
| 50 | 258 | 16.2 |
| 100 | 943 | 37.5 |
| 150 | 291 | 48.1 |
| 200 | 712 | 34 |
| 250 | 797 | 30.8 |
| 300 | 1555 | 46.3 |
| 350 | 508 | 42.3 |
| 400 | 419 | 28.6 |
| 450 | 463 | 39.5 |
| 500 | 688 | 28.6 |
| 550 | 879 | 12.9 |
| 600 | 554 | 19.2 |
| 650 | 1019 | 41.4 |
| 700 | 957 | 12.3 |
| 750 | 1883 | 12.2 |
| 800 | 3663 | 27.8 |
| 850 | 813 | 38.2 |
| 900 | 1521 | 27.4 |
| 950 | 996 | 37.8 |
| 1000 | 831 | 37.6 |
| 1500 | 3103 | 27.1 |
| 2000 | 2605 | 7.8 |
| 2500 | 2582 | 37.4 |
| 3000 | 3918 | 26.4 |
| 3500 | 3552 | 36.9 |
| 4000 | 3560 | 37 |
| 4500 | 2875 | 36.8 |
| 5000 | 6083 | 26.2 |
| 5500 | 3335 | 26.2 |
| 6000 | 3829 | 37.4 |
| 6500 | 3989 | 26.2 |
| 7000 | 5050 | 26.3 |
| 7500 | 8172 | 4.2 |
| 8000 | 6964 | 26.8 |
| 8500 | 13632 | 2.6 |
| 9000 | 41467 | 2.6 |
| 9500 | 12358 | 2.5 |
| 10000 | 4440 | 26.6 |
| 50000 | 7363 | 26.4 |
| 100000 | 97056 | 1.5 |

- Read


| Iterations | Read op rate(op/s) | Read latency mean |
|------------|--------------------|-------------------|
| 10 | 286 | 4.5 |
| 50 | 74 | 4.6 |
| 100 | 248 | 2.7 |
| 150 | 652 | 1.4 |
| 200 | 251 | 2.8 |
| 250 | 650 | 1 |
| 300 | 434 | 37.5 |
| 350 | 323 | 26.5 |
| 400 | 497 | 1 |
| 450 | 534 | 26.5 |
| 500 | 505 | 29 |
| 550 | 788 | 38 |
| 600 | 657 | 36.9 |
| 650 | 1118 | 0.8 |
| 700 | 896 | 26.1 |
| 750 | 7029 | 39.1 |
| 800 | 991 | 38.9 |
| 850 | 1839 | 26.8 |
| 900 | 1004 | 26.1 |
| 950 | 967 | 31.3 |
| 1000 | 1295 | 29 |
| 1500 | 2643 | 0.5 |
| 2000 | 2407 | 0.5 |
| 2500 | 14934 | 41.8 |
| 3000 | 3440 | 39.6 |
| 3500 | 11287 | 28.9 |
| 4000 | 2852 | 27.2 |
| 4500 | 7682 | 0.4 |
| 5000 | 15618 | 28.2 |
| 5500 | 7827 | 37.7 |
| 6000 | 15001 | 26.9 |
| 6500 | 9505 | 37.3 |
| 7000 | 15938 | 38.9 |
| 7500 | 12447 | 0.3 |
| 8000 | 10901 | 37.2 |
| 8500 | 12006 | 0.2 |
| 9000 | 4222 | 113.4 |
| 9500 | 10304 | 39.9 |
| 10000 | 11440 | 27.9 |
| 50000 | 14361 | 57.4 |
| 100000 | 14543 | 53.7 |

4.3. Graphs:

-Write



-Read



4.4. Observations:

As can be seen in the subsequent sections, the mean latency and op-rate for both read and write operations keep increasing, when increasing the number of iterations of read and write operations. The graph shows an exponential increase in, from the op-rates for both read and write operations.

4.5. Technical Challenges:

Major challenge in this topology was as was in topology 2 but after this docker couldn't be used.

Manually Cassandra was installed on each node

Configuration of Cassandra.yaml file was most challenging phase which consumed more than 20 hours to figure out and configure the node.

The cloud lab topology for the same is given in the below graph along with there IP's -



```
Topology View List View Manifest Graphs cn1 x cn2 x cn3 x
enp1s0f0 Link encap:Ethernet Hwaddr 70:e4:22:83:be:62
inet addr:128.104.222.202 Bcast:128.104.223.255 Mask:255.255.254.0
inet6 addr: fe80::72e4:22ff:fe83:be62/64 Scope:Link
UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
RX packets:62161572 errors:0 dropped:0 overruns:10 frame:0
TX packets:6203208 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1000
RX bytes:4820304052 (4.8 GB) TX bytes:954230861 (954.2 MB)
Memory:c6a00000-c6afffff

lo Link encap:Local Loopback
inet addr:127.0.0.1 Mask:255.0.0.0
inet6 addr: ::1/128 Scope:Host
UP LOOPBACK RUNNING MTU:65536 Metric:1
RX packets:477925 errors:0 dropped:0 overruns:0 frame:0
TX packets:477925 errors:0 dropped:0 overruns:0 carrier:0
collisions:0 txqueuelen:1
RX bytes:136255659 (136.2 MB) TX bytes:136255659 (136.2 MB)

dbms@cn1:~/apache-cassandra-3.11.2/bin$
```

```

Topology View List View Manifest Graphs cn1 cn2 cn3
dbms@cn2:~$ ifconfig
eno1 Link encap:Ethernet Hwaddr ec:bd:d7:85:5a:32
      inet addr:128.110.154.65 Bcast:128.110.155.255 Mask:255.255.252.0
      inet6 addr: fe80::eeb1:d7ff:fe85:5a32/64 Scope:Link
      UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
      RX packets:6123767 errors:0 dropped:0 overruns:0 frame:0
      TX packets:4921192 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1000
      RX bytes:896657490 (896.6 MB)  TX bytes:963581571 (963.5 MB)

lo Link encap:Local Loopback
      inet addr:127.0.0.1 Mask:255.0.0.0
      inet6 addr: ::1/128 Scope:Host
      UP LOOPBACK RUNNING  MTU:65536  Metric:1
      RX packets:110101 errors:0 dropped:0 overruns:0 frame:0
      TX packets:110101 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1
      RX bytes:11712760 (11.7 MB)  TX bytes:11712760 (11.7 MB)

dbms@cn2:~$

```

```

Topology View List View Manifest Graphs cn1 cn2 cn3
eno1 Link encap:Ethernet Hwaddr 34:17:eb:e5:72:6e
      inet addr:130.127.133.91 Bcast:130.127.135.255 Mask:255.255.252.0
      inet6 addr: fe80::3617:ebff:fee5:726e/64 Scope:Link
      UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
      RX packets:45646582 errors:0 dropped:0 overruns:0 frame:0
      TX packets:4492581 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1000
      RX bytes:3220895237 (3.2 GB)  TX bytes:856314843 (856.3 MB)
      Memory:91820000-9183ffff

lo Link encap:Local Loopback
      inet addr:127.0.0.1 Mask:255.0.0.0
      inet6 addr: ::1/128 Scope:Host
      UP LOOPBACK RUNNING  MTU:65536  Metric:1
      RX packets:106806 errors:0 dropped:0 overruns:0 frame:0
      TX packets:106806 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1
      RX bytes:10837841 (10.8 MB)  TX bytes:10837841 (10.8 MB)

dbms@cn3:~$

```

```

Topology View List View Manifest Graphs cn1 cn2 cn3
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Owns Host ID Rack
UN 128.110.222.202 414.71 KiB 256 ? 5fbc42b2-83c0-4373-a966-5957e34c0b57 222
Datacenter: 110
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Owns Host ID Rack
UN 128.110.154.65 412.34 KiB 256 ? 3845264d-5d42-4877-8554-ed65ea946640 154
Datacenter: 127
=====
Status=Up/Down
// State=Normal/Leaving/Joining/Moving
-- Address Load Tokens Owns Host ID Rack
UN 130.127.133.91 541.67 KiB 256 ? 33bf9d86-abdd-494e-abe2-42d3d047174a 133

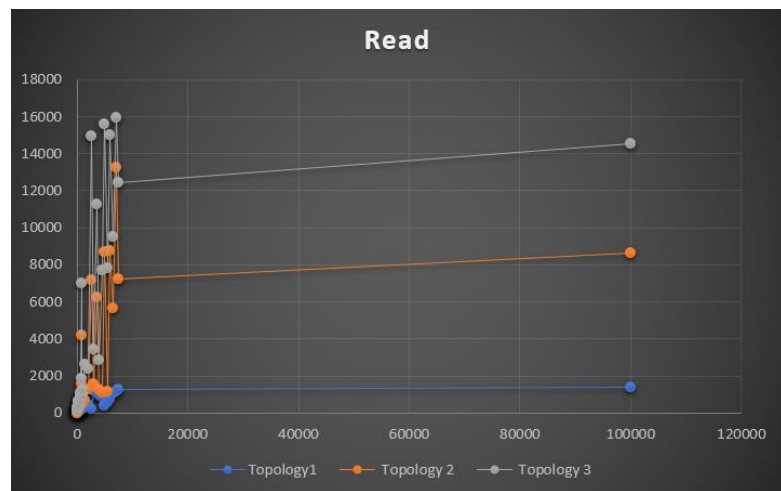
```

6. Conclusions:

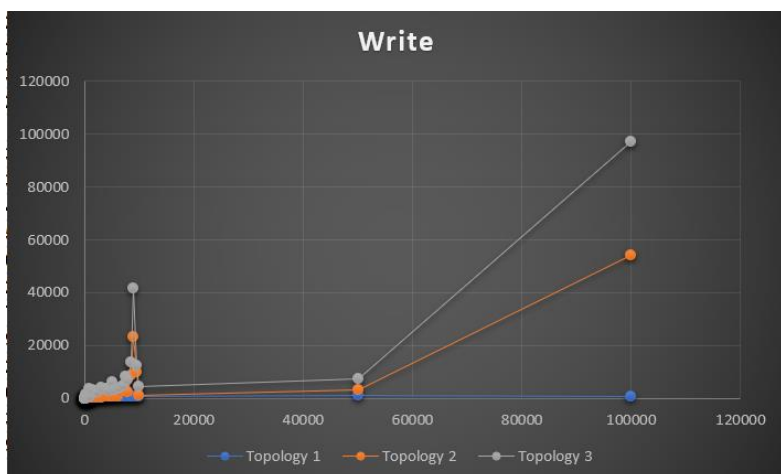
We can infer from the below graph that the latency will increase as the distance of the physical nodes in the topologies increases. Hence the curve of the topology one lies below in the graph where in topology 2 line

lies above the topology line1, but it comes below the topology 3 line. Hence in the distributed data bases we can say that latencies increase as the distance increases. This could be improved by using very good network infrastructure in the background. Future work for this project would be testing this scenario with that infrastructure.


-Read



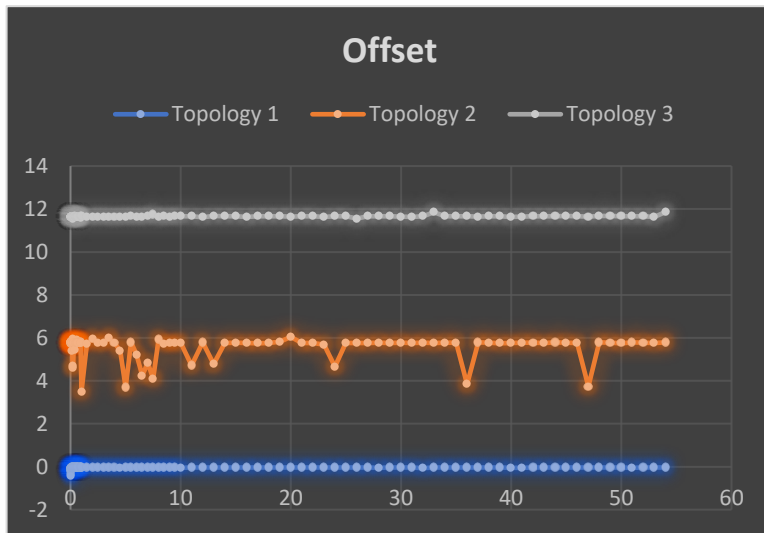
-Write



-Delay



-Offset



7. References

1. <https://hub.docker.com/r/bitnami/cassandra/>
2. <http://cassandra.apache.org/>
3. <https://help.ubuntu.com/lts/serverguide/NTP.html>
4. <https://blog.rapid7.com/2014/03/14/synchronizing-clocks-in-a-cassandra-cluster-pt-1-the-problem/>
5. <http://thelastpickle.com/blog/2017/02/08/Modeling-real-life-workloads-with-cassandra-stress.html>
6. <https://cloudlab.us/login.php>
7. <http://docs.cloudlab.us/users.html#%28part.join-project%29>

8. Links

All our work is saved and added in below Git repository.

GitHub:

https://github.com/dhawaskar/Cassandra_cluster_analysis